


# Data Analysis via Classification Algorithms

Lydia Y. Chen

Research Staff Member  
IBM Research Zurich Lab, Switzerland



# Classification Example: Email Spam

- Observe the set of emails (strings)
- Check if new emails are spam or non-spam

From: medshop@spam.com Subject: Viagra cheap meds...	Spam
From: my.instructor@Zurich.ibm.com Subject: important information here's how to ace the test...	No-spam
From: mike@example.org Subject: you need to see this how to win \$1,000,000...	???



# Classification Example: Credit Card Fraud

- Observe a set of old transactions
- Predict if new transactions are fraud or normal

Cash advance \$1,000 Cash advance \$10,000 Flight out of the country \$1,200	Fraud
Candy bar \$1.20 Groceries \$67.10 Restaurant \$35.82 ..... ...	Normal
Flight out of the country \$380 Hotel booking \$210 Groceries \$69.20	???



# Classification Example: Diabetes Diagnoses

- Observe a set of patients and their physical conditions
- Predict if new patients will have diabetes or not

Patient Alice Age 60 Blood pressure 150 mmHG Weight 78 kg  Patient Bob Age 30 Blood pressure 120 mmHG Weight 70 kg ..... ...	Diabetes     No Diabetes
Patient Maira Age 55 Blood pressure 150 mmHG Weight 60 kg	???



# Classification

# Classification Problem

- ▶ Input  $X$  (observable features)
  - ▶ One to multiple
  - ▶ E.g., age, blood pressure, and weight
- ▶ Output  $Y$  (class outcomes)
  - ▶ Binary ( $y=1$  or  $-1$ ) or multiple classes ( $G=1,2,..k$ )
  - ▶ E.g., the sign of diabetes or not
- ▶ Objective
  - ▶ Optimize the conditional probability of data set

$$\max Pr(G = k \mid X = x)$$

# Building Classifier

- ▶ Making assumptions about the data structure
- ▶ Choosing a classifier, e.g., linear v.s. nonlinear, generative v.s. discriminant
- ▶ Parameterizing/fitting the model (training phase), involving optimization procedures
- ▶ Classifying the new data (inference phase)

# Classification Algorithms

## Discriminative v.s. Generative

- ▶ Discriminative models capture  $P(Y | X)$  directly and generative models capture  $P(X), P(y)$
- ▶ Generative models are computational more expensive

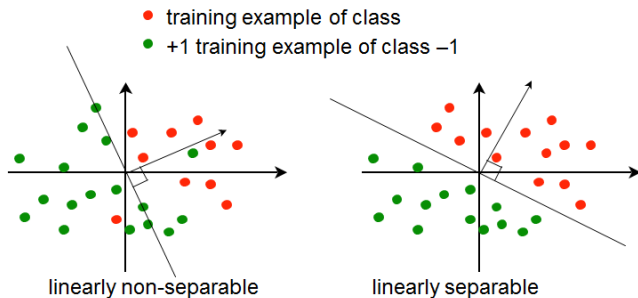


# Classification Algorithms

- ▶ Logistic regression (Discriminative)
- ▶ Decision Tree (Discriminative)
- ▶ Linear Discriminant analysis (Generative)
- ▶ Quadratic Discriminant analysis (Generative)
- ▶ Support Vector Machine (Discriminative)

# Linear Classification

- ▶ What is meant by linear classification?
  - ▶ The decision boundaries in the feature (input) space is linear
- ▶ May seem 'too simple', but in high-dimensions it is surprisingly powerful



# Linear Classification

- ▶ Training data  $(x_i, y_i)$  for  $i = 1 \dots N$ , and we wish to predict  $y'$  corresponding to new data vectors  $x'$
- ▶  $x_i \in R^n$ , and  $y_i \in \{0, 1\}$
- ▶ Predicted value  $\hat{y}$  is given by taking the sign of a linear function of the input:

$$\hat{y} = \text{sign}(\theta^T x), \text{sign}(z) = \begin{cases} 0 & \text{if } z \geq 0 \\ 1 & \text{if } z < 0 \end{cases}$$

# Logistic Regression

# Logistic Regression

- ▶ Classification based on probability
- ▶ Instead of just predicting the class, give the probability of the instance being that class
  - ▶ Directly learn conditional probability
$$Pr(G = k|X = x) = p_{\theta}(x)$$
- ▶ Focus on binary classification

## Simple Example

Cancer diagnosis from tumor size.



$$x = \begin{bmatrix} x_0 \\ x \end{bmatrix} = \begin{bmatrix} 1 \\ \textit{Tummor Size} \end{bmatrix}$$

$$\implies p(x) = 0.7$$

$\implies$  Patient that 70% chance of tumor being malignant

# Probability to Classification

- ▶ Logistic regression model assumes
  - ▶  $p_{\theta} = g(\theta^T x)$
  - ▶  $g(Z) = \frac{1}{1 + \exp(-Z)}$
- ▶ Assume a threshold and
  - ▶ Predict  $y = 1$ , if  $p_{\theta} \geq 0.5$
  - ▶ Predict  $y = 0$ , if  $p_{\theta} < 0.5$



# Assumptions to Model

- Logistic regression model

$$p_{\theta}(x) = \frac{1}{1 + \exp(-\theta^T x)}$$

$$\log \frac{Pr(y = 1 \mid x; \theta)}{Pr(y = 0 \mid x; \theta)} = \theta_0 + \theta_1 x_1 + \dots \theta_d x_d$$

- In other words, logistic regression assumes that the log odds is a linear function of

# Comparison with Linear Regression on Indicators

- ▶ Similarities:
  - ▶ Both attempt to estimate  $Pr(G = k \mid X = x)$
  - ▶ Both have linear classification boundaries
- ▶ Differences:
  - ▶ Linear regression on indicator matrix: approximate  $Pr(G = k \mid X = x)$  by a linear function of  $x$   
 $Pr(G = k \mid X = x)$  is not guaranteed to fall between 0 and 1 and to sum up to 1
  - ▶ Logistic regression:  $Pr(G = k \mid X = x)$  is a nonlinear function of  $x$ . It is guaranteed to range from 0 to 1 and to sum up to 1.



# Fitting Logistic Regression

- ▶ Criteria: find parameters that maximize the conditional likelihood of  $G$  given  $X$  using the training data.
  - ▶ Denote  $p_k(x_i; \theta) = \Pr(G = k | X = x_i; \theta)$ .
  - ▶ Given the first input  $x_1$ , the posterior probability of its class being  $g_1$  is  $\Pr(G = g_1 | X = x_1)$ .
- ▶ Since samples in the training data set are independent, the posterior probability for the  $N$  samples each having class  $g_i$ ,  $i = 1, 2, \dots, N$ , given their inputs  $x_1, x_2, \dots, x_N$  is:

$$\prod_{i=1}^N \Pr(G = g_i | X = x_i)$$

# Conditional Log Likelihood

$$\max \prod_{i=1}^N \Pr(G = g_i \mid X = x_i) = \max \log \prod_{i=1}^N \Pr(G = g_i \mid X = x_i)$$

- The conditional log-likelihood of the class labels in the training data set is

$$\begin{aligned} L(\theta) &= \sum_{i=1}^N \log \Pr(G = g_i \mid X = x_i) \\ &= \sum_{i=1}^N \log p_{g_i}(x_i; \theta) \end{aligned}$$

$i$  denotes the index of data

# Derivation of $L(\theta)$ for Binary Classifier

- ▶ For binary classification, if  $g_i = 1$ , denote  $y_i = 1$ ; if  $g_i = 2$ , denote  $y_i = 0$ .
- ▶ Let  $p_1(x; \theta) = p(x; \theta)$ , then  $p_2(x; \theta) = 1 - p_1(x; \theta) = 1 - p(x; \theta)$
- ▶ Since  $K = 2$ , the parameters  $\theta = \{\beta_{10}, \beta_1\}$ . We denote  $\beta = (\beta_{10}, \beta_1)^T$

## Derivation of $L(\theta)$ for Binary Classifier

- If  $y_i = 1$ , i.e.,  $g_i = 1$ ,

$$\begin{aligned}\log p_{g_i}(x; \beta) &= \log p_1(x; \beta) \\ &= 1 \cdot \log p_1(x; \beta) \\ &= y_i \cdot \log p_1(x; \beta)\end{aligned}$$

- If  $y_i = 0$ , i.e.,  $g_i = 2$ ,

$$\begin{aligned}\log p_{g_i}(x; \beta) &= \log p_2(x; \beta) \\ &= 1 \cdot \log (1 - p_1(x; \beta)) \\ &= (1 - y_i) \cdot \log (1 - p(x; \beta))\end{aligned}$$

- Since either  $y_i = 0$  or  $1 - y_i = 0$ , we have

$$\log p_{g_i}(x; \beta) = y_i \cdot \log p(x; \beta) + (1 - y_i) \cdot \log (1 - p(x; \beta))$$

# Binary Classifier of Logistic Regression

- Assumptions of Logistic Regression

$$p(x; \beta) = \Pr(G = 1 \mid X = x) = \frac{\exp(\beta^T x)}{1 + \exp(\beta^T x)}$$

$$1 - p(x; \beta) = \Pr(G = 2 \mid X = x) = \frac{1}{1 + \exp(\beta^T x)}$$

- The conditional likelihood

$$L(\beta) = \sum_{i=1}^N \log p_{g_i}(x_i; \beta)$$

- Put together

$$L(\beta) = \sum_i y_i \beta^T x_i - \log(1 + \exp \beta^T x_i)$$

# Finding $\beta$

How to solve

$$\max_{\beta} L(\beta) = \max_{\beta} \sum_i^N y_i \beta^T x_i - \log(1 + \exp \beta^T x_i)$$

- ▶ Non-linear optimization
- ▶ Stochastic gradient decent

## Example

- ▶ Diabetes data set. Input  $X$  is two dimensional.  $X_1$  and  $X_2$  are the two principal components of the original 8 variables.
- ▶ Class 1: without diabetes; Class 2: with diabetes.
- ▶ Applying logistic regression, we obtain

$$\beta = (0.7679, -0.6816, -0.3664)^T$$

## Example

- The posterior probabilities

$$Pr(G = 1 \mid X = x) = \frac{\exp(0.77 - 0.68x_1 - 0.37x_2)}{1 + \exp(0.77 - 0.68x_1 - 0.37x_2)}$$


$$Pr(G = 2 \mid X = x) = \frac{1}{1 + \exp(0.77 - 0.68x_1 - 0.37x_2)}$$

- The classification rule

$$G(\hat{x}) = \begin{cases} 0 & \text{if } 10.77 - 0.68x_1 - 0.37x_2 \geq 0 \\ 1 & \text{if } 10.77 - 0.68x_1 - 0.37x_2 < 0 \end{cases}$$



## Example



Solid line: decision boundary obtained by logistic regression.  
Dash line: decision boundary obtained by LDA (next topic).

# Discriminant Analysis

# Gaussian Discriminant Analysis

- ▶ Assume distribution of underlying data
- ▶ Robust for multi-classes analysis
- ▶ Classifiers can be obtained by closed form solutions
- ▶ And, we look into  $X = (X_1, X_2 \dots X_D)$

# Notion

- ▶ The prior probability of class  $k$  is  $\pi_k$ 
  - ▶  $\pi_k$  is usually estimated simply by empirical frequencies of the training set

$$\pi_k = \frac{\text{no. of samples in class } k}{\text{Total no. of samples}}$$

- ▶ The class-conditional density of  $X$  in class  $G = k$  is  $f_k(x)$
- ▶ Computing the posterior probability

$$Pr(G = k \mid X = x) = \frac{f_k(x)\pi_k}{\sum_{l=1}^K f_l(x)\pi_l}$$

- ▶ Obtain classifier By MAP (the Bayes rule for 0-1 loss)

$$\begin{aligned} \hat{G}(x) &= \arg \max_k Pr(G = k \mid X = x) \\ &= \arg \max_k f_k(x)\pi_k \end{aligned}$$

# Linear Discriminant Analysis

- ▶ Assumption 1: multivariate gaussian distribution




$$f_k(x) = \frac{1}{(2\pi)^{p/2}} \exp\left(-\frac{1}{2}(x - \mu_k)^T \sum_k^{-1} (x - \mu_k)\right)$$

E.g., Bivariate normal distribution



$$\mu_k = (0, 0), \text{ and } \Sigma_k = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

- ▶ Assumption 2:  $\Sigma_k = \Sigma, \forall k$ 
  - ▶ The Gaussian distributions are shifted versions of each other.

# Gaussian Distributions: Univariate



# Gaussian Distributions: Bivariate



# Optimal Classifier $G(\hat{x})$

$$\begin{aligned} G(\hat{x}) &= \arg \max_k Pr(G = k \mid X = x) \\ &= \arg \max_k f_k(x) \pi_k = \arg \max_k \log (f_k(x) \pi_k) \\ &= \arg \max_k [-\log((2\pi)^{p/2} |\Sigma|^{1/2}) \\ &\quad - \frac{1}{2}(x - \mu_k)^T \Sigma^{-1}(x - \mu_k) + \log(\pi_k)] \\ &= \arg \max_k -\frac{1}{2}(x - \mu_k)^T \Sigma^{-1}(x - \mu_k) + \log(\pi_k) \end{aligned}$$



# Optimal Classifier $G(\hat{x})$

- Define the linear discriminant function

$$G(\hat{x}) = \arg \max_k [x^T \Sigma^{-1} \mu_k - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log(\pi_k)]$$

- Then Classification rule

$$G(\hat{x}) = \arg \max_k \delta_k(x)$$

- The decision boundary between class  $k$  and  $l$  is:

$$\{x : \delta_k(x) = \delta_l(x)\}$$

- Or equivalently the following holds


$$\log \frac{\pi_k}{\pi_l} - \frac{1}{2} (\mu_k + \mu_l)^T \Sigma^{-1} (\mu_k - \mu_l) + x^T \Sigma^{-1} (\mu_k - \mu_l) = 0$$

# A Numerical Example

- ▶ Binary classification ( $k = 1, l = 2$ )
  - ▶ Define  $a_0 = \log \frac{\pi_1}{\pi_2} - \frac{1}{2}(\mu_1 + \mu_2)^T \Sigma^{-1}(\mu_1 - \mu_2)$
  - ▶ Define  $(a_1, a_2, \dots, a_D)^T = \Sigma^{-1}(\mu_1 - \mu_2)$
  - ▶ Classify to class 1 if  $a + \sum_{j=1}^D a_j x_k > 0$ ; to class 2 otherwise
- ▶ Example
  - ▶  $\pi_1 = \pi_2 = 0.5$ ,  $\mu_1 = (0, 0)^T$ , and  $\mu_2 = (0, 0)^T$
  - ▶  $\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 0.56 \end{bmatrix}$
  - ▶ Decision boundary

$$5.56 - 2.00x_1 + 3.56x_2 = 0$$

# Example



# Estimating Gaussian Distributions

- We need to estimate the Gaussian distribution

$$\hat{\pi}_k = N_k / N,$$

$$\hat{\mu}_k = \sum_{g_i=k} x^{(i)} / N_k$$

$$\hat{\Sigma} = \sum_{k=1}^K \sum_{g_i=k} (x^{(i)} - \hat{\mu}_k)(x^{(i)} - \hat{\mu}_k)^T / (N - K)$$

$N_k$  is the number of class-k samples

$x^{(i)}$  denotes the  $i^{(th)}$  sample vector

# Diabetes Data Set

- ▶ Two input variables computed from the principal components of the original 8 variables.

- ▶ Prior probabilities:  $\hat{\pi}_1 = 0.651$ ,  $\hat{\pi}_2 = 0.349$

$$\hat{\mu}_1 = (-0.43, -0.19)^T, \hat{\mu}_2 = (0.75, 0.36)^T$$


$$\Sigma = \begin{bmatrix} 1.79 & -0.15 \\ -0.15 & 1.66 \end{bmatrix}$$

- ▶ Classification rules:

$$G(\hat{x}) = \begin{cases} 1 & 0.77 - 0.67x_1 - 0.39x_2 \geq 0 \\ 2 & \text{otherwise} \end{cases}$$


## Example

The scatter plot follows. Without diabetes: stars (class 1), with diabetes: circles (class 2). Solid line: classification boundary obtained by LDA. Dash dot line: boundary obtained by linear regression of indicator matrix.



# Contour Plot

Contour plot for the density (mixture of two Gaussians) of the diabetes data




# Poor Results of LDA

- ▶ LDA is not necessarily bad when the assumptions about the density functions are violated.
- ▶ In certain cases, LDA may yield poor results.

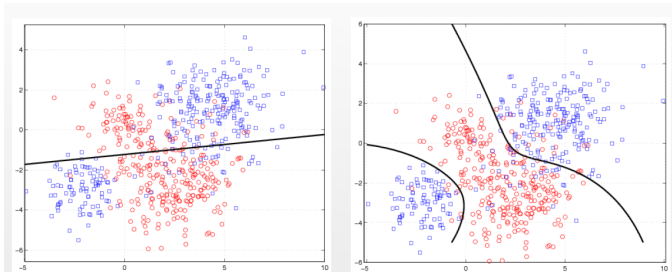


# Poor Results of LDA



Left: The true within class densities are Gaussian with identical covariance matrices across classes. Right: The true within class densities are mixtures of two Gaussians

# Poor Results of LDA



Left: Decision boundary by LDA. Right: Decision boundaries obtained by modeling each class by a mixture of two Gaussian

# Quadratic Discriminant analysis

- ▶ Estimate the covariance matrix  $\Sigma_k$  separately for each class  $k$ ,  $k = 1, 2, \dots, K$ .
- ▶ Quadratic discriminant function:

$$\delta_k(x) = \frac{1}{2} \log |\Sigma_k| - \frac{1}{2} (x - \mu_k)^T \Sigma_k^{-1} (x - \mu_k) + \log \pi_k$$

- ▶ Classification rule:


$$G(\hat{x}) = \arg \max_k \delta_k(x)$$

- ▶ Decision boundaries are quadratic equations in  $x$ .
- ▶ QDA fits the data better than LDA, but has more parameters to estimate.

# Diabetes Data Set

- ▶ Prior probabilities:  $\pi_1 = 0.651$ ,  $\pi_2 = 0.349$
- ▶ Prior probabilities:  $\hat{\pi}_1 = 0.651$ ,  $\hat{\pi}_2 = 0.349$
- ▶  $\hat{\mu}_1 = (-0.4035, -0.1935)^T$ ,  $\hat{\mu}_2 = (0.7528, 0.3611)^T$
- ▶  $\hat{\Sigma}_1 = \begin{bmatrix} 1.6769 & -0.1461 \\ -0.1461 & 1.5964 \end{bmatrix}$
- ▶  $\hat{\Sigma}_2 = \begin{bmatrix} 2.0087 & -0.3330 \\ -0.3330 & 1.7887 \end{bmatrix}$

## Example: QDA




# LDA on Expanded Bases

- ▶ Expand input space to include  $X_1, X_2, X_2$ , and  $X_2$ .
- ▶ Input is five dimensional:  $X = (X_1, X_2, X_1X_2, X_1^2, X_2^2)$ .
- ▶ Classification boundary:

$$0.65 - 0.73x_1 - 0.55x_2 - 0.006x_1x_2 - 0.07x_1^2 + 0.17x_2^2$$

- ▶ If the linear function on the right hand side is non-negative, classify as 1; otherwise 2

## Example: ELDA



Classification boundaries obtained by LDA using the expanded input space  $X_1, X_2, X_1X_2, X_1^2, X_1^2$ . Boundaries obtained by LDA and QDA

## Other Methods

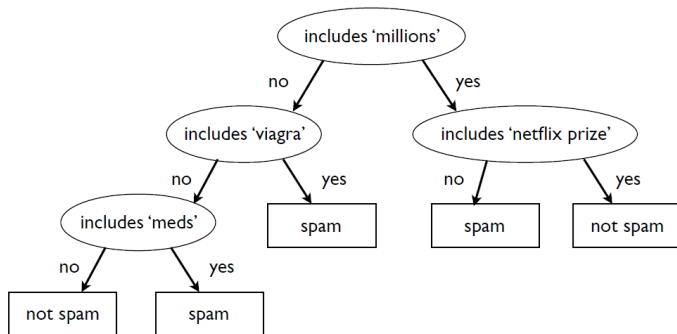


# Other Learning Approaches

- ▶ Rule based methods
  - ▶ Decision Tree
  - ▶ Random Forest
- ▶ Perceptron methods
  - ▶ Support Vector Machines
  - ▶ Neural Networks



# Decision Tree

- Idea: Ask a sequence of questions (as in the '20 questions' game) to infer the class





# Decision Tree: General Ideas I



- ▶ Simple idea: recursively divide up the space into pieces which are as pure as possible





# Decision Tree: General Ideas II



# Decision Tree: General Ideas III



# Decision Tree: General Ideas IV



# Decision Tree: General Ideas

- ▶ Key steps to build the tree
  1. A goodness of split criterion that can be evaluated for any split  $s$  of any node.
  2. A stop-splitting rule.
  3. A rule for assigning every terminal node to a class.
- ▶ How to measure if a subset of records is 'pure' or 'impure'
  - ▶ Entropy
  - ▶ Gini index
  - ▶ Classification error

# Characteristics of Decision Tree

- ▶ Nonparametric approach
  - ▶ Can approximate any decision boundary (hyper plane) to arbitrary precision
  - ▶ A practical starting point
- ▶ Local, greedy learning to find a reasonable solution in reasonable time
- ▶ Relatively easy to interpret (by experts or regular users of the system)
- ▶ Data fragmentation problem: the nodes far down the tree are based on a very small fraction of the data, even only on a few data points ) typically not very reliable information




# Support Vector Machine

- ▶ Construct linear decision boundaries that explicitly try to separate the data into different classes as well as possible.
- ▶ Good separation is defined in a certain form mathematically.
- ▶ Even when the training data can be perfectly separated by hyperplanes, LDA or other linear methods developed under a statistical framework may not achieve perfect separation.

# Optimal Separating Hyperplane

- The optimal separating hyperplane separates the two classes and maximizes the distance to the closest point from either class.



# Key Ideas of SVM

- ▶ Seek large margin separator to improve generalization
- ▶ Use optimization to find optimal hyperplane with few errors via slack variables

